

## 5.8 — Constexpr and consteval functions

by ALEX · JANUARY 19, 2024

In lesson [5.5 -- Constexpr variables](#) (<https://www.learncpp.com/cpp-tutorial/constexpr-variables/>), we introduced the `constexpr` keyword, which we used to create compile-time (symbolic) constants. We also introduced constant expressions, which are expressions that can be evaluated at compile-time rather than runtime.

Consider the following program, which uses two `constexpr` variables:

```
#include <iostream>

int main()
{
    constexpr int x{ 5 };
    constexpr int y{ 6 };

    std::cout << (x > y ? x : y) << " is greater!\n";

    return 0;
}
```

This produces the result:

```
6 is greater!
```

Because `x` and `y` are `constexpr`, the compiler can evaluate the constant expression `(x > y ? x : y)` at compile-time, reducing it to just `6`. Because this expression no longer needs to be evaluated at runtime, our program will run faster.

However, having a non-trivial expression in the middle of our print statement isn't ideal -- it would be better if the expression were a named function. Here's the same example using a function:

```
#include <iostream>

int greater(int x, int y)
{
    return (x > y ? x : y); // here's our expression
}

int main()
{
    constexpr int x{ 5 };
    constexpr int y{ 6 };

    std::cout << greater(x, y) << " is greater!\n"; // will be evaluated at runtime

    return 0;
}
```

This program produces the same output as the prior one. But there's a downside to putting our expression in a function: the call to `greater(x, y)` will execute at runtime. By using a function (which is good for modularity and documentation) we've lost our ability for that code to be evaluated at compile-time (which is bad for performance).



So how might we address this?

## Constexpr functions can be evaluated at compile-time

A **constexpr function** is a function whose return value may be computed at compile-time. To make a function a constexpr function, we simply use the `constexpr` keyword in front of the return type. Here's a similar program to the one above, using a constexpr function:

```
#include <iostream>

constexpr int greater(int x, int y) // now a constexpr function
{
    return (x > y ? x : y);
}

int main()
{
    constexpr int x{ 5 };
    constexpr int y{ 6 };

    // We'll explain why we use variable g here later in the lesson
    constexpr int g { greater(x, y) }; // will be evaluated at compile-time

    std::cout << g << " is greater!\n";

    return 0;
}
```

This produces the same output as the prior example, but the function call `greater(x, y)` will be evaluated at compile-time instead of runtime!

When a function call is evaluated at compile-time, the compiler will calculate the return value of the function call, and then replace the function call with the return value.



So in our example, the call to `greater(x, y)` will be replaced by the result of the function call, which is the integer value `6`. In other words, the compiler will compile this:

```
#include <iostream>

int main()
{
    constexpr int x{ 5 };
    constexpr int y{ 6 };

    constexpr int g { 6 }; // greater(x, y) evaluated and replaced with return value 6

    std::cout << g << " is greater!\n";

    return 0;
}
```

To be eligible for compile-time evaluation, a function must have a `constexpr` return type and not call any non-`constexpr` functions when evaluated at compile-time. Additionally, a call to the function must have arguments that are constant expressions (e.g. compile-time constant variables or literals).

## Author's note

We'll use the term "eligible for compile-time evaluation" later in the article, so remember this definition.

## For advanced readers

There are some other lesser encountered criteria as well. These can be found [here](https://en.cppreference.com/w/cpp/language/constexpr) (<https://en.cppreference.com/w/cpp/language/constexpr>).

Our `greater()` function definition and function call in the above example meets these requirements, so it is eligible for compile-time evaluation.

## Constexpr functions can also be evaluated at runtime

Functions with a constexpr return value can also be evaluated at runtime, in which case they will return a non-constexpr result. For example:

...



```
#include <iostream>

constexpr int greater(int x, int y)
{
    return (x > y ? x : y);
}

int main()
{
    int x{ 5 }; // not constexpr
    int y{ 6 }; // not constexpr

    std::cout << greater(x, y) << " is greater!\n"; // will be evaluated at runtime

    return 0;
}
```

In this example, because arguments `x` and `y` are not constant expressions, the function cannot be resolved at compile-time. However, the function will still be resolved at runtime, returning the expected value as a non-constexpr `int`.

## Key insight

Allowing functions with a constexpr return type to be evaluated at either compile-time or runtime was allowed so that a single function can serve both cases.

Otherwise, you'd need to have separate functions (a function with a constexpr return type, and a function with a non-constexpr return type). This would not only require duplicate code, the two functions would also need to have different names!

This is also why C++ does not allow constexpr function parameters. A constexpr function parameter would imply the function could only be called with a constexpr argument. But this is not the case -- constexpr functions can be called with non-constexpr arguments when the function is evaluated at runtime.

## So when is a constexpr function evaluated at compile-time?

You might think that a constexpr function would evaluate at compile-time whenever possible, but unfortunately this is not the case.

According to the C++ standard, a constexpr function that is eligible for compile-time evaluation must be evaluated at compile-time if the return value is used where a constant expression is required. Otherwise, the compiler is free to evaluate the function at either compile-time or runtime.

Let's examine a few cases to explore this further:

• • •



```
#include <iostream>

constexpr int greater(int x, int y)
{
    return (x > y ? x : y);
}

int main()
{
    constexpr int g { greater(5, 6) };           // case 1: always evaluated at compile-time
    std::cout << g << " is greater!\n";

    std::cout << greater(5, 6) << " is greater!\n"; // case 2: may be evaluated at either runtime or compile-time

    int x{ 5 }; // not constexpr but value is known at compile-time
    std::cout << greater(x, 6) << " is greater!\n"; // case 3: likely evaluated at runtime

    std::cin >> x;
    std::cout << greater(x, 6) << " is greater!\n"; // case 4: always evaluated at runtime

    return 0;
}
```

In case 1, we're calling `greater()` with constant expression arguments, so it is eligible to be evaluated at compile-time. The initializer of `constexpr` variable `g` must be a constant expression, so the return value is used in a context that requires a constant expression. Thus, `greater()` must be evaluated at compile-time.

In case 2, the `greater()` function is again being called with constant expression arguments, so it is eligible for compile-time evaluation. However, the return value is not being used in a context that requires a constant expression (`operator<<` always executes at runtime), so the compiler is free to choose whether this call to `greater()` will be evaluated at compile-time or runtime!

In case 3, we're calling `greater()` with one argument that is not a constant expression. However, this argument has a value that is known at compile-time. Under the as-if rule, the compiler could decide to treat `x` as `constexpr`, and evaluate this call to `greater()` at compile-time. But more likely, it will evaluate it at runtime.

In case 4, the value of argument `x` can't be known at compile-time, so this call to `greater()` will always evaluate at runtime.

Note that your compiler's optimization level setting may have an impact on whether it decides to evaluate a function at compile-time or runtime. This also means that your compiler may make different choices for debug vs. release builds (as debug builds typically have optimizations turned off).

• • •



## Key insight

A `constexpr` function must be evaluated at compile-time if the return value is used where a constant expression is required. Otherwise,

compile-time evaluation is not guaranteed.

Thus, a `constexpr` function is better thought of as “can be used in a constant expression”, not “will be evaluated at compile-time”.

## Key insight

Put another way, we can categorize the likelihood that a `constexpr` function will actually be evaluated at compile-time as follows:

Always (required by the standard):

- `constexpr` function is called in context where constant expression is required.

Probably (there's little reason not to):

- `constexpr` function is called in context where constant expression isn't required, all arguments are constant expressions.

Possibly (if optimized under the as-if rule):

- `constexpr` function is called in context where constant expression isn't required, some arguments are not constant expressions but their values are known at compile-time.

Never (not possible):

- `constexpr` function is called in context where constant expression isn't required, some arguments are not constant expressions and their values are not known at compile-time.

## Determining if a `constexpr` function call is evaluating at compile-time or runtime

Prior to C++20, there are no standard language tools available to do this.

In C++20, `std::is_constant_evaluated()` (defined in the `<type_traits>` header) returns a `bool` indicating whether the current function call is executing in a constant context. This can be combined with a conditional statement to allow a function to behave differently when evaluated at compile-time vs runtime.

```
#include <type_traits> // for std::is_constant_evaluated

constexpr int someFunction()
{
    if (std::is_constant_evaluated()) // if compile-time evaluation
        // do something
    else // runtime evaluation
        // do something else
}
```

Used cleverly, you can have your function produce some observable difference (such as returning a special value) when evaluated at compile-time, and then infer how it evaluated from that result.

## Forcing a `constexpr` function to be evaluated at compile-time



There is no way to tell the compiler that a `constexpr` function should prefer to evaluate at compile-time whenever it can (even in cases where the return value is used in a non-constant expression).

However, we can force a `constexpr` function that is eligible to be evaluated at compile-time to actually evaluate at compile-time by ensuring the return value is used where a constant expression is required. This needs to be done on a per-call basis.

The most common way to do this is to use the return value to initialize a `constexpr` variable (this is why we've been using variable 'g' in prior examples). Unfortunately, this requires introducing a new variable into our program just to ensure compile-time evaluation, which is ugly and

## For advanced readers

There are several hacky ways that people have tried to work around the problem of having to introduce a new `constexpr` variable each time we want to force compile-time evaluation. See [here](https://quuxplusone.github.io/blog/2018/08/07/force-constexpr/) (<https://quuxplusone.github.io/blog/2018/08/07/force-constexpr/>) and [here](https://artificial-mind.net/blog/2020/11/14/cpp17-consteval/) (<https://artificial-mind.net/blog/2020/11/14/cpp17-consteval/>).

However, in C++20, there is a better workaround to this issue, which we'll present in a moment.

## Consteval C++20

C++20 introduces the keyword **consteval**, which is used to indicate that a function must evaluate at compile-time, otherwise a compile error will result. Such functions are called **immediate functions**.

...



```
#include <iostream>

consteval int greater(int x, int y) // function is now consteval
{
    return (x > y ? x : y);
}

int main()
{
    constexpr int g { greater(5, 6) };           // ok: will evaluate at compile-time
    std::cout << g << '\n';

    std::cout << greater(5, 6) << " is greater!\n"; // ok: will evaluate at compile-time

    int x{ 5 }; // not constexpr
    std::cout << greater(x, 6) << " is greater!\n"; // error: consteval functions must evaluate at compile-time

    return 0;
}
```

In the above example, the first two calls to `greater()` will evaluate at compile-time. The call to `greater(x, 6)` cannot be evaluated at compile-time, so a compile error will result.

## Best practice

Use `consteval` if you have a function that must run at compile-time for some reason (e.g. performance).

## Using `consteval` to make `constexpr` execute at compile-time C++20

The downside of `consteval` functions is that such functions can't evaluate at runtime, making them less flexible than `constexpr` functions, which can do either. Therefore, it would still be useful to have a convenient way to force `constexpr` functions to evaluate at compile-time (even when the return value is being used where a constant expression is not required), so that we could have compile-time evaluation when possible, and runtime evaluation when we can't.

`Consteval` functions provides a way to make this happen, using a neat helper function:

```

#include <iostream>

// Uses abbreviated function template (C++20) and `auto` return type to make this function work with any type of value
// See 'related content' box below for more info (you don't need to know how these work to use this function)
constexpr auto compileTime(auto value)
{
    return value;
}

constexpr int greater(int x, int y) // function is constexpr
{
    return (x > y ? x : y);
}

int main()
{
    std::cout << greater(5, 6) << '\n';           // may or may not execute at compile-time
    std::cout << compileTime(greater(5, 6)) << '\n'; // will execute at compile-time

    int x { 5 };
    std::cout << greater(x, 6) << '\n';           // we can still call the constexpr version at runtime if we wish

    return 0;
}

```

This works because `constexpr` functions require constant expressions as arguments -- therefore, if we use the return value of a `constexpr` function as an argument to a `constexpr` function, the `constexpr` function must be evaluated at compile-time! The `constexpr` function just returns this argument as its own return value, so the caller can still use it.

Note that the `constexpr` function returns by value. While this might be inefficient to do at runtime (if the value was some type that is expensive to copy, e.g. `std::string`), in a compile-time context, it doesn't matter because the entire call to the `constexpr` function will simply be replaced with the calculated return value.

## Related content

We cover `auto` return types in lesson [10.9 -- Type deduction for functions](https://www.learncpp.com/cpp-tutorial/type-deduction-for-functions/) (<https://www.learncpp.com/cpp-tutorial/type-deduction-for-functions/>). We cover abbreviated function templates (`auto` parameters) in lesson [11.8 -- Function templates with multiple template types](https://www.learncpp.com/cpp-tutorial/function-templates-with-multiple-template-types/) (<https://www.learncpp.com/cpp-tutorial/function-templates-with-multiple-template-types/>).

## Constexpr/constexpr functions are implicitly inline

Because `constexpr` functions may be evaluated at compile-time, the compiler must be able to see the full definition of the `constexpr` function at all points where the function is called. A forward declaration will not suffice, even if the actual function definition appears later in the same compilation unit.

This means that a `constexpr` function called in multiple files needs to have its definition included into each such file -- which would normally be a violation of the one-definition rule. To avoid such problems, `constexpr` functions are implicitly inline, which makes them exempt from the one-definition rule.

As a result, `constexpr` functions are often defined in header files, so they can be #included into any .cpp file that requires the full definition.

`constexpr` functions are also implicitly inline (presumably for consistency).

## Rule

The compiler must be able to see the full definition of a `constexpr` or `constexpr` function, not just a forward declaration.

## Best practice

`constexpr/constexpr` functions used in a single source file (.cpp) can be defined in the source file above where they are used.

`constexpr/constexpr` functions used in multiple source files should be defined in a header file so they can be included into each source file.

## Constexpr/constexpr function parameters are not constexpr (but can be used as arguments to other constexpr functions)

The parameters of a `constexpr` function are not `constexpr` (and thus cannot be used in constant expressions). Such parameters can be declared as `const` (in which case they are treated as runtime constants), but not `constexpr`. This is because a `constexpr` function can be evaluated at runtime (which wouldn't be possible if the parameters were compile-time constants).

However, an exception is made in one case: a `constexpr` function can pass those parameters as arguments to another `constexpr` function, and that subsequent `constexpr` function can be resolved at compile-time. This allows `constexpr` functions to still be resolved at compile-time when

they call other `constexpr` functions (including themselves recursively).

Perhaps surprisingly, the parameters of a `constexpr` function are not considered to be `constexpr` within the function either (even though `constexpr` functions can only be evaluated at compile-time). This decision was made for the sake of consistency.

Here's a contrived example illustrating this:

```
#include <iostream>

constexpr int goo(int c)
{
    return c;
}

constexpr int foo(int b)
{
    constexpr int b2 { b }; // compile error: b is not a constant expression within foo()

    return goo(b);          // okay: b can still be used as argument to constexpr function goo()
}

int main()
{
    constexpr int a { 5 };

    std::cout << foo(a); // okay: constant expression a can be used as argument to constexpr function foo()

    return 0;
}
```

## Related content

If you need `constexpr` parameters to a function (e.g. to use somewhere that requires a constant expression), see [11.9 -- Non-type template parameters](#) (<https://www.learncpp.com/cpp-tutorial/non-type-template-parameters/>).

## Can a `constexpr` function call a non-`constexpr` function?

The answer is yes, but only when the `constexpr` function is being evaluated in a non-constant context. A non-`constexpr` function may not be called when a `constexpr` function is evaluating in a constant context (because then the `constexpr` function wouldn't be able to produce a compile-time constant value).

Calling a non-`constexpr` function is allowed so that a `constexpr` function can do something like this:

```
#include <type_traits> // for std::is_constant_evaluated

constexpr int someFunction()
{
    if (std::is_constant_evaluated()) // if compile-time evaluation
        return someConstexprFcn();    // calculate some value at compile time
    else
        // runtime evaluation
        return someNonConstexprFcn(); // calculate some value at runtime
}
```

Now consider this variant:

```
constexpr int someFunction(bool b)
{
    if (b)
        return someConstexprFcn();
    else
        return someNonConstexprFcn();
}
```

This is legal as long as `someFunction(false)` is never called in a constant expression.

The C++ standard says that a `constexpr` function must return a `constexpr` value for at least one set of arguments, otherwise it is technically ill-formed. Therefore, calling a non-`constexpr` function unconditionally in a `constexpr` function makes the `constexpr` function ill-formed. However, compilers are not required to generate errors or warnings for such cases -- therefore, the compiler probably won't complain unless you try to call such a `constexpr` function in a constant context.

Therefore, we'd advise the following:

1. For best results, avoid calling non-`constexpr` functions from within a `constexpr` function if possible.
2. If your `constexpr` function requires different behavior for constant and non-constant contexts, conditionalize the behavior with `if (std::is_constant_evaluated())`.

3. Always test your `constexpr` functions in a constant context, as they may work when called in a non-constant context but fail in a constant context.

## Why not `constexpr` every eligible function?

There are a few reasons you may not want to `constexpr` a function:

1. `constexpr` is part of the interface of a function. Once a function is made `constexpr`, it can be called by other `constexpr` functions or used in contexts that require constant expressions. Removing the `constexpr` later will break such code.
2. `constexpr` makes functions harder to debug because you can't inspect them at runtime.

But as a general rule, if you can `constexpr` a function, you should.

### Best practice

Unless you have a specific reason not to, a function that can be made `constexpr` generally should be made `constexpr`.



[Next lesson](#)

5.9 [Introduction to std::string](#)



[Back to table of contents](#)



[Previous lesson](#)

5.7 [Inline functions and variables](#)

Leave a comment...

Name\*

Notify me about replies:

[POST COMMENT](#)

Email\*



Find a mistake? Leave a comment above! [?](#)

Avatars from <https://gravatar.com/> are connected to your provided email address.

145 COMMENTS

Newest ▾

We and our partners share information on your use of this website to help improve your experience.

Do not sell my info:



OKAY