

10.5 — Arithmetic conversions

 ALEX  JANUARY 13, 2024

In lesson [6.1 -- Operator precedence and associativity](https://www.learncpp.com/cpp-tutorial/operator-precedence-and-associativity/) (<https://www.learncpp.com/cpp-tutorial/operator-precedence-and-associativity/>), we discussed how expressions are evaluated according to the precedence and associativity of their operators.

Consider the following expression:

```
int x { 2 + 3 };
```

When binary operator+ is invoked, it is given two operands, both of type `int`. Because both operands are of the same type, that type will be used to perform the calculation and to return the result. Thus, `2 + 3` will evaluate to `int` value `5`.

But what happens when the operands of a binary operator are of different types?

```
??? y { 2 + 3.5 };
```

In this case, operator+ is being given one operand of type `int` and another of type `double`. Should the result of the operator be returned as an `int`, a `double`, or possibly something else altogether? When defining a variable, we can choose what type it has. In other cases, for example when using `std::cout <<`, the type the calculation evaluates to changes the behavior of what is output.

In C++, certain operators require that their operands be of the same type. If one of these operators is invoked with operands of different types, one or both of the operands will be implicitly converted to matching types using a set of rules called the **usual arithmetic conversions**.

The operators that require operands of the same type

The following operators require their operands to be of the same type:

- The binary arithmetic operators: `+`, `-`, `*`, `/`, `%`
- The binary relational operators: `<`, `>`, `<=`, `>=`, `==`, `!=`
- The binary bitwise arithmetic operators: `&`, `^`, `|`
- The conditional operator `?:` (excluding the condition, which is expected to be of type `bool`)

The usual arithmetic conversion rules

The usual arithmetic conversion rules are somewhat complex, so we'll simplify a bit. The compiler has a ranked list of types that looks something like this:

- long double (highest rank)
- double
- float
- long long
- long
- int (lowest rank)

The following rules are applied to find a matching type:

- If one operand is an integral type and the other a floating point type, the integral operand is converted to the type of the floating point operand (no integral promotion takes place).

Otherwise, any integral operands are numerically promoted (see [10.2 -- Floating-point and integral promotion](https://www.learncpp.com/cpp-tutorial/floating-point-and-integral-promotion/) (<https://www.learncpp.com/cpp-tutorial/floating-point-and-integral-promotion/>)).

- If one operand is signed and the other unsigned, special rules apply (see below)
- Otherwise, the operand with lower rank is converted to the type of the operand with higher rank.

For advanced readers

The special matching rules for integral operands with different signs:

- If the rank of the unsigned operand is greater than the rank of the signed operand, the signed operand is converted to the type of the unsigned operand.
- If the type of the signed operand can represent all the values of the type of the unsigned operand, the type of the unsigned operand is converted to the type of the signed operand.
- Otherwise both operands are converted to the corresponding unsigned type of the signed operand.

Related content

You can find the full rules for the usual arithmetic conversions [here](https://en.cppreference.com/w/cpp/language/usual_arithmetic_conversions) (https://en.cppreference.com/w/cpp/language/usual_arithmetic_conversions).

Some examples

Binary operator+ is one of the operators that requires its operands to have the same type.

In the following examples, we'll use the `typeid` operator (included in the `<typeinfo>` header), to show the resulting type of an expression.

First, let's add an `int` and a `double`:

```
#include <iostream>
#include <typeinfo> // for typeid()

int main()
{
    int i{ 2 };
    double d{ 3.5 };
    std::cout << typeid(i + d).name() << ' ' << i + d << '\n'; // show us the type of i + d

    return 0;
}
```

In this case, the `double` operand has the highest priority, so the lower priority operand (of type `int`) is type converted to `double` value `2.0`. Then `double` values `2.0` and `3.5` are added to produce `double` result `5.5`.

On the author's machine, this prints:

```
double 5.5
```

Note that your compiler may display something slightly different, as the output of `typeid.name()` is left up to the compiler.

Now let's add two values of type `short`:

```
#include <iostream>
#include <typeinfo> // for typeid()

int main()
{
    short a{ 4 };
    short b{ 5 };
    std::cout << typeid(a + b).name() << ' ' << a + b << '\n'; // show us the type of a + b

    return 0;
}
```

Because neither operand appears on the priority list, both operands undergo integral promotion to type `int`. The result of adding two `ints` is an `int`, as you would expect:

```
int 9
```

Signed and unsigned issues

This prioritization hierarchy and conversion rules can cause some problematic issues when mixing signed and unsigned values. For example, take a look at the following code:

```
#include <iostream>
#include <typeinfo> // for typeid()

int main()
{
    std::cout << typeid(5u-10).name() << ' ' << 5u - 10 << '\n'; // 5u means treat 5 as an unsigned integer

    return 0;
}
```

You might expect the expression `5u - 10` to evaluate to `-5` since `5 - 10 = -5`. But here's what actually results:

```
unsigned int 4294967291
```

Due to the conversion rules, the `int` operand is converted to an `unsigned int`. And since the value `-5` is out of range of an `unsigned int`, we get a result we don't expect.

Here's another example showing a counterintuitive result:

```
#include <iostream>

int main()
{
    std::cout << std::boolalpha << (-3 < 5u) << '\n';

    return 0;
}
```

While it's clear to us that `5` is greater than `-3`, when this expression evaluates, `-3` is converted to a large `unsigned int` that is larger than `5`. Thus, the above prints `false` rather than the expected result of `true`.

This is one of the primary reasons to avoid unsigned integers -- when you mix them with signed integers in arithmetic expressions, you're at risk for unexpected results. And the compiler probably won't even issue a warning.



[Next lesson](#)

10.6 [Explicit type conversion \(casting\) and static_cast](#)



[Back to table of contents](#)



[Previous lesson](#)

10.4 [Narrowing conversions, list initialization, and constexpr initializers](#)

Leave a comment...



Name*



Email*



Notify me about replies:



POST COMMENT



Find a mistake? Leave a comment above!



Avatars from <https://gravatar.com/> are connected to your provided email address.

We and our partners share information on your use of this website to help improve your experience.



Do not sell my info: ☐

OKAY